

Prompt Syntax (PS), Specification Draft v0.2

Status: Working draft, design decisions marked  (settled),  (proposed, awaiting confirmation),  (open question)
Editors: [author], with drafting assistance **Last updated:** 2026-08-02

Candidate change for v0.2.2: proposes **standing promotion** (§13.2): an application may declare a versioned, origin-bound authoring surface through which a named agent principal can request a closed set of actions. The declaration changes which segments are eligible for parsing; it does not widen the capability envelope. Each exercise remains subject to current authorization and produces a typed fill outcome in the trace. Motivation: v0.2.1 required authority-checked promotion but illustrated only per-event user acceptance, leaving no specified reusable form for agentic applications. Source: an author-controlled application review (2026-08-01) that found a live capability-escalation instance of the gap. Also proposed: explicit producer and framing bounds, replay handling, closed-action confusable handling, and open question 16.

Changelog v0.1 → v0.2: major revision incorporating external platform-operator review (2026-07-19). Added: authority/specificity precedence split with non-escalation rule (§10); four-layer architecture (§4); trace visibility tiers (§12); pinned canonical identifiers (§6.2); capability negotiation (§7); parameter normalization records (§7.2); standardized failure codes (§8.2); budget construct replacing generic `limit()` (§9.1); fallback authorization rule (§9.2); determinism taxonomy (§14); mandatory island highlighting for interactive implementations (§13). Deferred: multi-model span execution (→ orchestration extension). Reframed: portability as *portable intent + transparent resolution, not portable execution*.

1. Abstract

Prompt Syntax (PS) is a vendor-neutral language for expressing **authorized prompt-execution intent**, together with a conformance protocol describing how that intent was resolved and fulfilled. Prompts remain natural language with embedded *syntax islands*, directives that give users **routing control** (deterministic, *authority-bounded* selection of models, tools, skills, and context) and **routing transparency** (an inspectable record of assembly, resolution, and fulfillment: the **Prompt Trace**). PS defines a grammar, resolution semantics, an authority model, and a conformance contract; it does not define presentation UX, and it promises portable *expression of intent*, never portable execution.

2. Terminology

Term	Definition
Directive	A PS construct embedded in a prompt that requests assembly, routing, or execution behavior.
Syntax island	A directive occurrence inside otherwise-plain natural language.
Entity	A referenceable resource: model, agent, plugin/connector, tool, skill, file, context source.
Canonical identifier	The immutable, versioned, vendor-qualified form of an entity reference (§6.2).
Resolver	The defined procedure mapping an entity reference to a canonical identifier.
Authority level	The governance tier a configuration value originates from (§10.1).
Capability envelope	The set of permissions and resources granted by all higher authority levels; directives operate only within it.

Term	Definition
Fulfillment policy	The directive's contract for unsatisfiable requirements: strict / partial / best-effort (§8).
Fill report	The per-requirement requested → normalized → applied record, with reason codes (§8.2).
Budget	A venue-enforced execution bound with defined measurement points (§9.1).
Fallback chain	An ordered list of routing steps tried on fill failure; the only contingency construct in PS (§9.2).
Canonical form	The fully explicit representation every PS prompt compiles to (§11).
Prompt Trace	The tiered, machine-readable artifact recording assembly, resolution, and fulfillment (§12).
Trace tier	A defined visibility level of the trace: user / developer / operator / oversight (§12.2).
Routing control	The property that directives <i>guarantee</i> what runs, within the capability envelope, rather than the model inferring intent.
Routing transparency	The property that the user can inspect, at their tier, what will run (pre-execution) and what ran (post-execution).
Promotion	An authority-checked change in provenance type that makes a specific segment eligible for parsing. Promotion never grants a capability.
Declared authoring surface	A versioned, origin-bound interface through which output from a named principal may be parsed for a closed set of actions (§13.2).
PS-shaped	Descriptive only: an implementation uses PS syntax or one or more PS-aligned behaviors but does not satisfy a named layer-and-version conformance claim.

3. Design principles

1. **Graceful degradation.** A prompt with zero directives is valid PS; every directive is optional prose-compatible.
2. **Familiarity borrowing.** Every sigil reuses a convention with an existing large user base (@-mentions, /commands, XML tags, frontmatter).
3. **Progressive disclosure.** Casual form → parameterized form → strict form, all compiling to one canonical representation; the equivalence is a property of the grammar.
4. **Content safety.** Pasted code/HTML/XML must never accidentally become directives; only reserved namespaces are live.
5. **Contract, not presentation.** The spec mandates information availability (the Prompt Trace); presentation is implementation-defined, with one exception (§13, island highlighting).
6. **Minimal reserved surface.** One live tag namespace (ps), two sigils (@, /), one fence (--ps). The reserved set never grows across versions.
7. **One reference grammar.** The @ entity-reference syntax is the sole routing syntax, reused verbatim at every scope.
8. **Non-escalation.**  (v0.2) Prompt content can *narrow* the capability envelope; it can never expand it. No directive grants access, permission, or budget that higher authority has not already granted.

9. **Intent portability, not execution portability.** ✓ (v0.2) PS guarantees that intent is expressed portably and resolution is reported transparently; it never guarantees that execution is identical across venues.

4. Layered architecture ✓ (v0.2)

PS is specified as four layers with independent conformance. A venue may claim conformance per layer; layers evolve at different rates (the grammar is stable for years; parameters and runtimes are not).

Layer	Contents	Stability
PS/Core	Grammar, escaping, scopes, canonical references, round-trip property	High
PS/Capabilities	Entity discovery, capability documents, parameter registries, authorization surface	Medium
PS/Execution	Fulfillment policies, budgets, fallback chains	Medium
PS/Trace	Trace schema, tiers, fill reports, failure codes	Medium-high

A conformance claim names layers and version: e.g., PS/Core 0.2 + PS/Trace 0.2 (user tier). An implementation that borrows PS syntax or behavior without meeting such a claim is **PS-shaped** only. PS-shaped is not a conformance level and MUST NOT be presented as conformance. The adoption path is PS-shaped, then named-layer conformance, with independent certification left to the governance work in OQ#15.

5. Directive scopes ✓

Scope	Syntax	Applies to	Example
Point	@ entity reference; / action	The whole message (routing) or the reference site (context)	@opus(temperature: 0.2) Summarize this. / concise
Span	Reserved tag <ps ...></ps>	The enclosed text region	<ps @file:glossary.md>use these terms</ps>
Document	Frontmatter fence ---ps ... ---	The entire prompt; the strict form	see §11

5.1 Span tag forms ✓

The ps tag accepts the same reference grammar as point scope (principle 7): <ps @ref ...> casual; <ps key="value"> canonical/XML-friendly (compiler-emitted).

5.2 v0.1/v0.2 span restriction ✓ (v0.2, was open; settled per external review)

Spans are restricted to: **context inclusion, local instruction scope, tool availability, visibility/provenance marking**. Multi-model span execution (<ps @haiku>...</ps> routing a region to a different model) is **deferred to a separate orchestration extension**: it requires execution-order, data-flow, context-isolation, token-accounting, failure- and injection-boundary semantics that do not belong in the core language. The syntax remains forward-compatible; venues MUST reject (not ignore) model references in spans with SPAN_EXECUTION_UNSUPPORTED.

6. Entity references and resolution

6.1 Reference syntax

```
reference ::= "@" [ namespace ":" ] name [ "!" ] [ arglist ]
namespace ::= "model" | "agent" | "plugin" | "tool" | "skill" | "file" | "ctx" | vendor-ext
name ::= identifier ( "." | "/" ) identifier )* [ "@" version ]
arglist ::= "(" arg ( "," arg )* ")"
arg ::= key ":" value (* JSON5-style scalars *)
action ::= "/" name [ arglist ]
```

6.2 Canonical identifiers (v0.2)

Casual aliases (@opus, @fast) are legal input. Resolution MUST bind every reference to a **globally canonical, immutable, versioned identifier**, recorded in the canonical form and the trace:

```
@opus           → @model:anthropic/claude-opus-4-5@2026-05-01
@tool:search-contracts → @tool:workspace/acme/search-contracts@3
@file:q3-report.md   → @file:workspace/acme/9fc2...
```

Model aliases drift; dated snapshots do not. The trace never contains an unpinned binding.

6.3 Ambiguity and namespace resolution

1. A bare name matching exactly one entity in the capability envelope binds to it.
2. Multiple matches are an **ambiguity condition**: implementations MUST surface it and MUST NOT silently guess (ENTITY_AMBIGUOUS).
3. Every binding, including the deciding rule, is recorded in the trace's resolution report.

? Open: default namespace search order (proposal: session bindings → project config → model → agent → plugin → tool → skill → file).

6.4 Environment-relative parsing, documented limitation (v0.2)

Because bare-name recognition depends on installed entities, identical text may parse differently across venues. PS guarantees a **visible** parse, not an identical parse. Normative mitigations: island highlighting is MUST for interactive implementations (§13); **consequential actions** (tool execution, connector access, data egress) MUST be either fully qualified or explicitly confirmed by the user; venues SHOULD offer a strict mode in which only <ps> / ---ps content is executable.

7. Capabilities and parameters (PS/Capabilities)

7.1 Capability document (v0.2)

A conformant venue publishes a machine-readable capability document; implementations negotiate against it before fill:

```
ps_version: "0.2"
supports:
  scopes: [point, span, document]
  fulfillment: [strict, partial, best-effort]
  budgets: [estimated_cost, observed_cost, wall_time, first_token]
  cancellation: preferred # required | preferred | unsupported
  trace_tiers: [user, developer]
namespaces:
  model: { discovery: "...", pinning: date-snapshot }
  tool: { discovery: "...", }
```

A prompt MAY declare the PS version it targets (ps: "0.2" in frontmatter); venues MUST reject versions they cannot honor rather than degrade silently.

7.2 Parameters: normalization, not passthrough (v0.2, replaces v0.1 passthrough)

There is no stable universal parameter registry. PS defines a small **normalized core** (temperature, top_p, max_tokens, stop, seed, effort) plus vendor-native names. Every parameter's journey is recorded as a **normalization record** in the fill report:

```
{ "requested": { "effort": "max" },
  "normalized": { "reasoning.effort": "xhigh" },
  "applied": { "reasoning.effort": "xhigh" },
  "status": "translated",
  "mapping_rule": "vendor-profile/openai-2026-07" }
```

Status is one of applied | translated | clamped | unsupported | ignored, each with a defined fill consequence under the directive's policy. Undocumented passthrough is non-conformant: it creates the appearance of control without semantics.

8. Fulfillment (PS/Execution)

8.1 Policies (settled; external review concurs)

Inspired by exchange order types: a routing directive is an order; the venue fills it or reports why not.

Policy	Trading analog	Semantics
strict	Fill-or-Kill	Every requirement satisfied exactly, or the directive fails pre-execution. No silent substitution.
partial	Immediate-or-Cancel	Fill what can be filled; every drop recorded.
best-effort	Market order	Venue discretion; all degradations still recorded.

Default : **split**, strict on entity binding (@opus means opus, never a silent substitute), partial on parameters (unsupported values drop with a fill-report entry). ! **semantics**  **v0.2.1, clarified per external review**: since entity binding is already strict by default, ! escalates the **entire directive** to strict: every parameter becomes a hard requirement (no drop, no clamp; translation only where the vendor profile declares it semantics-preserving), and any degradation is a fill failure instead of a reported drop. @opus = strict entity, partial params; @opus! = strict everything. Named forms: fill="strict" (span/document), fill: strict (frontmatter; document level = all-or-none). ? *Open (severity/consent): degradation-severity classes and an optional confirm-on-degrade preflight mode for consequential settings, a post-hoc report is not consent.*

8.2 Failure codes (v0.2)

Fill failures and degradations carry standardized codes (extensible registry, PS_-free core set):

```
ENTITY_NOT_FOUND      ENTITY_AMBIGUOUS      ENTITY_UNAUTHORIZED
PARAM_UNSUPPORTED     PARAM_CLAMPED         PARAM_IGNORED
BUDGET_PREFLIGHT_REJECTED BUDGET_EXCEEDED      TIMEOUT_NOT_CANCELLED
DATA_POLICY_BLOCKED   SPAN_EXECUTION_UNSUPPORTED VERSION_UNSUPPORTED
FALLBACK_EXHAUSTED
```

9. Budgets and fallback chains (PS/Execution)

9.1 Budgets (v0.2, replaces generic limit())

limit(...) remains as syntax sugar, expanding to a **budget** with defined measurement points:

```
budget:
  estimated_cost_max: $0.10 # pre-flight: reject if floor estimate exceeds
  observed_cost_max:  $0.15 # runtime: abort (if cancellable) when exceeded
  wall_time_max:      5s    # submission → completion, tool calls included
  first_token_max:    1s    # streaming responsiveness bound
  cancellation:       preferred # what the directive requires: required | preferred
  scope:              attempt # attempt | chain
```

Semantics: a violated budget is a fill failure. Venues declare their enforcement capability in the capability document (§7.1); a directive requiring cancellation: required fails pre-flight (BUDGET_PREFLIGHT_REJECTED) against a venue declaring

unsupported. A timeout the venue could not cancel is reported as `TIMEOUT_NOT_CANCELLED` with observed cost, enforcement is honest-by-declaration, never pretended. Implementations **MUST NOT** silently degrade output (e.g., truncate) to satisfy a budget. scope: chain bounds the **total** across fallback attempts (*resolves v0.1 open question*).

9.2 Fallback chains (structure settled; authorization rule added in v0.2)

PS has **no general conditionals**, no predicates over content, no branching on output (delegated to orchestration frameworks). The only re-routing trigger is a **fill failure**. Precedent: CSS font stacks; trading's contingent orders.

```
route    ::= step ( "else" step ) * [ "else" terminal ]
step     ::= reference [ limitclause ]
terminal ::= "fail" | "ask"

@opus!(temperature: 0.2) limit(wall_time: 5s, estimated_cost: $0.10)
  else @model:openai/gpt-5.6@2026-07-01 limit(wall_time: 3s)
  else ask
```

First successful fill wins; the fill report records every attempt with outcome, codes, and measured cost/latency. **Authorization rule** ( **v0.2**): every step is checked against the same capability envelope, data classification, consent, allowlists, residency, retention, connector authorization. A failed primary route never makes a prohibited secondary route permissible; a cross-provider fallback is a data-governance decision, not merely an availability decision, and blocked steps report `DATA_POLICY_BLOCKED`.

10. Authority and precedence (v0.2, replaces v0.1 single chain; critical fix)

Precedence has **two independent dimensions**. The v0.1 single chain (`system < ... < span`) was unsafe: read literally, user content could override platform policy.

10.1 Authority precedence (who may decide)

```
platform policy > organization/workspace policy > application/developer policy
                > user policy > prompt content
```

Lower authority may **narrow** permissions, budgets, and entity sets; it may never expand them (principle 8).

10.2 Specificity precedence (which value wins, within one authority level)

```
document < point < span          (inner/more-specific wins)
```

10.3 Normative rule

A more specific directive overrides configuration **only within the capability envelope granted by higher-authority policy**. An out-of-envelope directive is a fill failure (`ENTITY_UNAUTHORIZED` / `DATA_POLICY_BLOCKED`), never a silent ignore, the user always learns that authority, not availability, was the reason.

Every envelope constraint that suppressed a directive is recorded in the trace at the requesting user's tier (as a typed refusal, not necessarily with the policy's content).

10.5 Refusal reasons and verified attributes (v0.2.1)

A directive can be blocked for categorically different reasons that current interfaces collapse into one undifferentiated outcome (e.g., a single "switched model" banner covering safety blocks, capacity substitution, and org-policy denials alike). PS requires the reason be **typed** in the trace, distinguishing at minimum: availability (entity unavailable/capacity), **safety** (`SAFETY_BLOCKED`, with a category and the provenance of the *triggering segment*, which may be a retrieved document or tool result, not the user's message), org/data policy (`DATA_POLICY_BLOCKED`), and authorization (`ENTITY_UNAUTHORIZED`). The receipt names the reason, the deciding authority, and, normatively, an available **recourse path** where one exists (appeal, verification program, feedback channel).

Verified principal attributes (the credential-informed envelope). A safety or policy envelope **MAY** be widened for a principal presenting a *verified attribute* (an identity proof, a credential such as a cyber-verification-program membership, an institutional affiliation). This is expressed entirely within the authority model and is **not** non-escalation-violating: the wider envelope is granted by a *higher* authority (the verifier/provider reading the attribute), never claimed by prompt content. The signed attribute is an **input to** the authority-holder's decision, never a self-granted capability; the decision and the

attributes it relied on (by reference, not by disclosing the credential) are recorded. Verified-attribute policies MUST NOT be mandated by PS, identity-gating carries chilling, privacy, and gatekeeping costs, and whether to use the mechanism is contestable policy left to the authority. (*Existing practice: providers already run verification programs that widen safeguards for verified defensive-security researchers, PS standardizes the seam, not the program.*)

? *Open (companion-paper territory): appeal-loop semantics; credential formats and privacy-preserving attribute proofs; interaction of verified envelopes with the oversight tier.*

10.4 Policy system integration (PEP/PDP) ♦ (v0.2.1)

PS does not define a policy language, and it adds **no new policy decisions**, the organization keeps its existing engine. What PS adds is what the existing control stack cannot produce: (1) **prompt-level decision semantics**, network- and key-layer controls are blind inside the request; PS's decision requests carry the entities, classifications, and fallback routes a meaningful policy needs to see; and (2) **per-request enforcement evidence**, today an admin setting is a request to the vendor with no proof it fired; typed refusals and fill reports make every policy outcome verifiable. PS defines the **enforcement seam**: the PS runtime is a *policy enforcement point* (PEP); the organization's policy engine is the *policy decision point* (PDP), attached at the org/workspace authority level. At resolution time, for each directive, the PEP submits a decision request, **principal** (identity/groups from the IdP), **entity reference** (canonical identifier), **context labels** (data classification of referenced files/sources), **jurisdiction/residency**, and the PDP answers: in-envelope or refused (with code), applicable budget, permitted egress set for fallback steps. Decisions and refusals are recorded in the trace (policy hits, not policy content). **Bindings** to concrete engines (OPA/Rego, Cedar, XACML-style) are external versioned documents, declared in the capability document.

Worked example (♦ v0.2.1, illustrative; decision-interface schema remains OQ#10). User input: @opus! Summarize @file:q3-acquisition-memo.pdf else @gemini-flash. At resolution time, before assembly, before any call, the PEP submits per-binding decision requests to the org's existing engine:

```
{ "principal": { "id": "user:meas@acme", "groups": ["finance"] },
  "action": "route",
  "route": [
    { "step": 1, "entity": "model:anthropic/claude-opus-4-5@2026-05-01",
      "provider_jurisdiction": "EU" },
    { "step": 2, "entity": "model:google/gemini-flash@2026-06-01",
      "provider_jurisdiction": "US", "role": "fallback" } ],
  "context": [ { "entity": "file:acme/9fc2...", "classification": "confidential" } ],
  "budget_scope": "attempt" }
```

Against an unchanged org policy (e.g., Rego/Cedar: *confidential sources route only to EU-approved models*), the PDP returns per-step decisions: step 1 permit, step 2 deny (egress). The PEP fills step 1 and marks step 2 DATA_POLICY_BLOCKED **before any failure could trigger it**, the chain is judged as one decision, not discovered call-by-call, and the trace records the policy_decision_ids.

Why a call-level gateway PDP is not equivalent. Gateways with policy sidecars can and do filter API calls today. But a call-level decision point sees flattened text and a model name; it cannot see (1) **typed references before flattening**, by the time the gateway inspects the call, the confidential file is extracted text and its classification label died in assembly, whereas PS decides on the reference while the label exists; (2) **the chain**, the fallback call does not exist yet at gateway time, and when it fires it is a fresh call that may individually pass; (3) **injected expansion**, which skill added what, or prompt-content attempts at tool-access escalation (inert-content rule, §13); (4) **its own effectiveness**, a gateway drop is silent; a PS refusal is typed, user-visible, and receipted. The layers compose: gateways remain useful as defense in depth below the PS seam.

Note, the externalization pattern: vendor profiles (§7.2), oversight profiles (§12.3), and policy bindings (§10.4) are the same architectural move: the core spec defines interface + semantics; stakeholder-specific mappings live outside and version independently. Enterprise-relevant consequences already in core: typed refusals as policy audit events; org-authority budgets; the fallback-authorization/egress rule (§9.2); pinned identifiers as model change management (§6.2); capability documents as procurement artifacts (§7.1).

11. Canonical form and the strict document

Canonical form: frontmatter (all routing/config, canonical pinned identifiers, defaults materialized, PS version declared) + body (prose with islands replaced by resolved references).

```
---ps
version: "0.2"
route:
  - { model: "model:anthropic/claude-opus-4-5@2026-05-01",
      params: { temperature: 0.2 }, fill: strict,
      budget: { wall_time_max: 5s, estimated_cost_max: $0.10 } }
  - { model: "model:openai/gpt-5.6@2026-07-01" }
skills: [ "skill:workspace/acme/concise@1" ]
context: [ "file:workspace/acme/9fc2..." ]
---
```

Summarize the report.

Round-trip property (normative): `compile(casual) == compile(strict(compile(casual)))`. The intent-reification layer **MUST** output this strict form, editable and re-submittable.

12. The Prompt Trace (PS/Trace)

12.1 Structure (settled; normative schema: `prompt-trace.schema.json`)

The trace is a **turn-level object** with an ordered events array interleaving *inference records* and *boundary events*. Three settled decisions:

- **Shape (D1): flat segments + provenance steps.** Ordered segment list whose concatenation reproduces the compiled prompt byte-exactly within tier scope (rule R1); each injected segment carries a pointer into a step table whose parent chains reconstruct the expansion tree (R2, R3). Source-map / W3C-PROV pattern.
- **Turn scope (D2): tiered, with the boundary rule.** Routing record + fill report **MUST** for every venue invocation; full segments **MUST** for user-initiated inferences, **SHOULD** for internal continuations (absence *declared* via coverage, never silent, R6). Everything crossing the venue boundary (tool I/O) **MUST** be recorded; tool internals are explicitly outside the perimeter.
- **Content (D3): hybrid three-state union.** `inline` ($\leq$ threshold; floor 4 KB for the user tier, R4) | `external` (sha256 + length + handle; dedups across loop steps) | `withheld` (salted digest + length band + minimum-visible tier). Oversight-tier exports **MUST** support **materialized mode**, external content resolved and embedded at export time, with blob retention governed at oversight scope (R5).

MUST: compiled prompt (tier-scoped); resolution report (every binding + deciding rule, ambiguities, envelope refusals); routing record incl. **fill report** (§8.2 codes, normalization records, per-attempt outcomes with measured cost/latency). **SHOULD:** dry-run (pre-execution) trace; stable segment IDs for diffing. **MAY:** any presentation.

12.2 Visibility tiers (v0.2, replaces *flat redaction*)

Tier	Audience	Contents
User	End user	User-visible sources, selected tools, applied parameters, typed envelope refusals, fill report
Developer	Application owner	1. application instructions, tool definitions, runtime events
Operator	Platform personnel	1. infrastructure and policy diagnostics
Oversight	Auditors & regulators (via profiles, §12.3)	Evidence-grade export per the applicable oversight profile

Completeness is **tier-relative**: a tier's trace is complete with respect to content at or below that tier's authority; higher-tier content appears as a **typed, sized placeholder** (`{type: system_policy, tier: operator}`). ? Open ("the transparency floor"): what placeholder detail is safe, exact lengths and content hashes can leak (dictionary attacks against candidate texts); banding and salted digests are candidate mitigations. This is a negotiation between auditability and confidentiality, flagged as a research question. Trace availability interacts with venue data-retention configuration; the capability document declares supported tiers.

12.3 Oversight: evidence substrate and profiles ✓ (v0.2.1)

Enterprise audit and government/regulatory oversight are **distinct lenses**. An auditor's access is granted by contract *inside* the authority ladder; a regulator's access is compelled by law from **above platform policy** (§10 note below). Regulatory requirements are jurisdictional, sectoral, and mutually conflicting (e.g., EU AI Act Art. 12 automatic record-keeping vs. GDPR minimization vs. WORM-retention financial rules), and they demand an evidentiary standard, tamper-evidence, authenticity, chain of custody (contributing to, not constituting, legal non-repudiation), that internal audit does not.

PS therefore specifies no jurisdiction's rules. Instead:

- **Evidence substrate (PS/Trace core)**. Properties any oversight regime composes from: integrity chaining or signing of trace records (SHOULD in core; oversight profiles are expected to make these MUST, with canonical serialization, key custody, and verification procedures profile-specified); completeness attestations (what coverage the venue claims, signed); **signed sequence numbers with explicit gap records** (an omitted event leaves a visible hole, not silence); trusted timestamps (SHOULD); retention and residency metadata fields (MUST when the oversight tier is supported); **provable redaction**, a withheld segment's existence, kind, and timing are demonstrable without revealing content; the fully **materialized export** (§ Decision 3) with blob retention governed at oversight scope. Because a venue-produced trace is self-asserted, profiles MAY additionally require independent trust mechanisms: append-only transparency logs, cross-reconciliation against gateway or billing records, independently operated collectors, or remote attestation.
- **Oversight profiles (external documents)**. Jurisdiction- or framework-specific mappings (e.g., an EU-AI-Act profile, a SOC 2 profile) declaring which substrate properties are MUST and with what parameters (retention period, residency, signature regime). Profiles are versioned and maintained outside the core spec, the PDF/A / FHIR national-profile / OTel semantic-convention pattern. The capability document declares supported profiles.

Authority note (amends §10.1): the authority ladder is crowned by an acknowledged but externally governed level, *law and regulation*, above platform policy. Non-escalation continues to bind prompt content; disclosure obligations, however, may be compelled from this level, making the transparency floor (§12.2) **authority-relative**: content a venue may withhold at the user tier can be compellable at the oversight tier under an applicable profile.

13. Content safety and input marking ✓

- **Provenance-typed parsing (inert content) ✓ (v0.2.1, security-critical, from external review)**: Only explicitly designated **authoring segments**, text the principal typed or an authority-checked surface produced as PS, are parsed for directives. Retrieved documents, file contents, tool results, quoted messages, and model-generated text are typed **inert**: directive-shaped text inside them (e.g., a document containing `@tool:send_email!(...)`) is content, never a directive, and cannot introduce directives except through an explicit, authority-checked *promotion* operation recorded in the trace (per-event acceptance, or the proposed declared standing surface in §13.2). This rule is prior to and stronger than fence exemption; without it, PS would be a prompt-injection and capability-escalation surface.
- **The parameterization principle ✓ (v0.2.1)**: security rests on provenance, never on escaping. PS is parsed only at the interface layer, only over authored segments, models never execute PS, so inert content is never "escaped for safety"; it is never parsed at all. Content inclusion is **by reference** (canonical id + digest), the discipline that ended SQL injection: data is never concatenated into the command channel. Escapes (`\@`, `\`, fences) exist solely as authoring ergonomics for deliberate literals; no security property may depend on them.
- **Inner content and nesting ✓ (v0.2.1)**: a span's inner prose is authored text, islands within it parse normally; content *resolved into* a span via its references remains inert. Canonical-form serialization MUST carry content segments by reference or in typed data fields, never spliced into directive text, so serialize-then-recompile cannot promote inert content (round-trip safety). Nested-span conflict semantics: pending conflict table (OQ). **Model-generated PS** (including the reification layer's proposed strict form) is inert until promoted: explicitly accepted by the user (per-event reification

acceptance, the archetypal promotion), or, where supported, exercised through a declared standing surface (§13.2). Either way the promotion is authority-checked and recorded in the trace.

- Only the ps tag namespace is live; all other <...> is content.
- @// islands require a word boundary and a resolvable or qualified name; emails and paths do not parse.
- Escapes \@, \/, and fenced code blocks are always content.
- **Island highlighting is MUST for interactive implementations** (v0.2, upgraded from *SHOULD per review*): every recognized directive is visually marked before submission. Non-interactive (API) callers use dry-run traces for the same guarantee.

13.1 Internationalization ♦ (v0.2.1)

PS prompts are natural language in any script; the grammar must not assume Latin text.

- **Sigils are script-neutral, with width equivalence.** @ and / are language-independent and present on keyboards worldwide; parsers MUST treat full-width forms (U+FF20 @, U+FF0F /), commonly produced by CJK input methods, as equivalent to their ASCII forms.
- **Unicode identifiers.** Entity names follow Unicode identifier rules (UAX #31) with NFC normalization, @file:អ្នកបោះពុម្ព .md is valid. Namespace keywords, policy names, and else remain English in the **canonical form** (the portability lesson of Excel’s localized function names: localized *canonical* syntax destroys portability). Implementations MAY accept localized input aliases for keywords and units; these compile to canonical form, internationalization rides the same casual→strict machinery as everything else, and the round-trip property applies.
- **Script-aware boundaries.** The “word boundary before a sigil” rule is defined by Unicode segmentation (UAX #29), not ASCII \b, scripts without inter-word spaces (Khmer, Thai, Japanese, Chinese) MUST still get correct island detection. Conformance test corpora MUST include non-spacing-script cases.
- **Bidirectional text safety.** In RTL contexts (Arabic, Hebrew), editors SHOULD wrap islands in bidi isolates (FSI...PDI) for stable rendering. Security-normative: parsers MUST reject or neutralize bidi control characters *inside* directive islands, Trojan Source-style reordering (CVE-2021-42574) could otherwise visually disguise what a directive does, defeating mandatory highlighting.
- **Confusables.** Homoglyph references (@opus with Cyrillic o) are a spoofing vector against entity binding. Open-name-space resolvers SHOULD apply confusable detection (UTS #39) and surface skeleton matches as ambiguity conditions rather than silently binding. For small, closed identifier sets such as declared action names, canonical identifiers SHOULD be ASCII-restricted and implementations MAY accept ASCII case-folded aliases. A non-ASCII confusable does not bind as an alias and, once it appears in an authoring segment, produces a typed failure. Prose arguments remain Unicode content.
- **Localized scalars.** Budget/parameter input MAY use locale conventions (decimal comma, local currency); canonical form and traces always use canonical units (USD, ms/s, dot-decimal).

? *Open: governance of localized keyword-alias tables (per-language registries? vendor-defined?); which locales the conformance suite must cover.*

13.2 Standing promotion: declared authoring surfaces ♦ (proposed for v0.2.2)

PS governs the forward channel, principal to venue. Agentic applications also run a continuous **reverse channel**: the application acts on model output to write records, open pull requests, or edit files. Section 13 types model-generated text inert. v0.2.1 required any promotion to be explicit, authority-checked, and traced, but illustrated only per-event user acceptance. It did not specify a reusable promotion form for a machine principal.

A standing declaration fills that gap without weakening provenance safety. It changes only whether a segment is eligible for parsing. It does **not** widen an authority or capability envelope, and it does not make the segment trustworthy. Resolution, policy, and authorization are evaluated again for every requested action.

Proposed normative rule. An application MAY, at application/developer authority or above, publish a versioned declaration that output from a named agent principal through a named, origin-bound frame is a **declared authoring surface**. The declaration is standing (made once, not accepted by a user on every event) and MUST specify:

1. **Identity and lifecycle.** A stable declaration identifier and version, the declaring authority, the producer principal or principal class, the scope in which it is valid, and how it is activated, revoked, and replaced. Model text cannot create, amend, or reactivate its own declaration.
2. **Origin and framing.** The transport, message or invocation boundary, output field, and exact segment framing. Hosts MUST preserve those boundaries and parse the designated frame directly. Searching an undifferentiated transcript,

- concatenated prose, quoted content, or retrieved data for a delimiter is NOT a valid surface. A delimiter found inside inert content cannot create an authoring segment.
3. **Live namespace and action set.** The vendor-ext namespace whose references are live inside the frame, plus a closed set of actions and parameter schemas. One surface being live grants no meaning to any other namespace, entity, tool, or action. Actions MAY be authority-differentiated, and an agent-authored segment MUST NOT claim user provenance or user authority.
 4. **Capability and authorization bound.** The surface operates strictly within the current envelope of the declared producer and application. Non-escalation (principle 8) applies unchanged. Cross-scope reach MUST be explicit, and every exercise MUST be authorized against the envelope and policy in force at execution time. A stale declaration is never an authorization grant.
 5. **Fill and replay behavior.** Every qualified or otherwise directive-shaped reference inside the frame ends in exactly one recorded outcome. A reference that fails to parse, bind, or authorize is a typed, surfaced failure, never a silent no-op. A frame is consumed at most once unless the declaration defines an idempotent replay rule; rejected duplicates also receive a typed outcome. Alias acceptance is MAY; accepted aliases compile to canonical form.
 6. **Trace obligations.** The declaration is recorded in or referenced by the capability document (§7.1). Each exercise records the declaration id and version, declaring and producing principals, invocation and frame identifiers, provenance type, current authorization decision, replay disposition, and per-reference fill outcomes. Segment content follows the tier-scoped inline | external | withheld rules of §12.1 rather than being copied into every trace tier.

Content outside a declared frame remains inert per §13. Reification acceptance is the per-event, user-authority form of promotion; a declared authoring surface is a reusable, higher-authority provenance rule. Both make a segment parse-eligible, and neither grants capability by itself.

Security boundary. A standing declaration authorizes a machine principal to request bounded actions. It does not establish that the principal’s output is correct, that the principal was not influenced by indirect prompt injection, or that the requested action reflects user intent. The control limits effects, fails closed outside the declared surface, and makes each outcome auditable. Consequential actions remain subject to the qualification or confirmation requirement in §6.4.

Confusables in closed action sets. Declared actions follow the closed-identifier guidance in §13.1. Inside a declared frame, a rejected confusable MUST produce a typed failure rather than becoming inert or binding silently.

14. Determinism taxonomy (v0.2)

Documentation and conformance language MUST distinguish, and never conflate:

Kind	PS guarantees it?
Routing-contract determinism , parsing, canonicalization, resolution, and policy evaluation under a pinned environment; execution choices recorded	Yes (within envelope; this is PS’s claim, renamed from “execution determinism” per external review: tool calls, races, and provider behavior are not deterministic and are <i>recorded</i> , not guaranteed)
Generation determinism , identical output text	No (even with seed; vendor implementation varies)
Semantic reproducibility , same meaning across runs/venues	No (retrieval, web content, tools, policies drift)

@model:x(seed: 42)! guarantees the seed was applied, not that the output is identical.

15. Out of scope (delegated)

Output/decoding constraints (LMQL, Guidance), params/normalization is the integration point. Agent orchestration graphs (DSPy, LangGraph) and multi-model span execution, orchestration extension. Content-conditional branching, orchestration frameworks.

16. Open questions queue

1. Default namespace search order (§6.3).
2. ~~Trace JSON schema decisions~~ **Resolved**, flat+provenance / tiered turn scope / hybrid content confirmed; normative schema published (§12.1).

3. The transparency floor: placeholder detail vs. leakage (§12.2).
4. `else ask` terminal: interaction contract with intent reification.
5. Capability-document schema detail (§7.1) and vendor-profile registry governance.
6. Failure-code registry governance (who mints codes; vendor extensions).
7. Vendor mapping table (Claude, ChatGPT, Cursor, Copilot, Gemini), separate document.
8. Orchestration extension scoping (multi-model spans, §5.2).
9. Oversight-profile governance: who authors and maintains jurisdiction profiles (standards body? consortium?); how profile conformance is asserted and verified (§12.3).
10. PDP decision-interface schema (§10.4): request/response shape, latency budget for per-directive policy evaluation, caching semantics, and binding governance.
11. (*from external auditor-lens review, 2026-07-19*) Identifier governance: who controls canonical-identifier uniqueness, vendor prefixes, deprecation, alias ownership, disputes (§6.2).
12. Budget measurement profiles: tariff/price-sheet version identifiers, cached-token treatment, queueing inclusion, cancellation-attempt evidence (§9.1).
13. Trace-privacy deployment profiles: field-level encryption, access logging, deletion, legal hold, retention conflicts, purpose limitation (§12.2).
14. Version-baseline consolidation: merge v0.2.1 markers into a clean v0.3 baseline before public spec release; conformance tests reference exactly one baseline.
15. Conformance vs. certification: negative/adversarial test requirements, signed capability documents, and a path to independent verification (currently: two author-controlled implementations demonstrate spec precision, not third-party assurance).
16. (*from author-controlled agentic-application review, 2026-08-01*) Standing-promotion governance (§13.2): declaration schema detail in the capability document and Prompt Trace; whether routing directives (model or tool references) are ever legal inside an agent-authored surface; nested interactions between surfaces and spans; origin-bound framing across streaming transports; replay and idempotency profiles; and migration guidance for PS-shaped implementations toward honest named-layer conformance claims.

*Naming: properties are routing **control** and routing **transparency**; the artifact is the **Prompt Trace**; fulfillment vocabulary is strict/partial/best-effort.*